



The **AiEdge** ACADEMY

Fine-Tuning LLMs




The Different Fine-Tuning Tasks

- For Language modeling
- For Sentence prediction
- For Text classification
- For Token classification
- For Text encoding
- For Multimodal modeling




Causal Language Modeling






Causal Language Modeling






Causal Language Modeling






Causal Language Modeling






Causal Language Modeling






Causal Language Modeling

Causal Mask:
attention computed
only for previous
words






Self-attentions






Masked Language Modeling




Causal Language Modeling



Masked Language Modeling




Masked Language Modeling






Masked Language Modeling






Masked Language Modeling



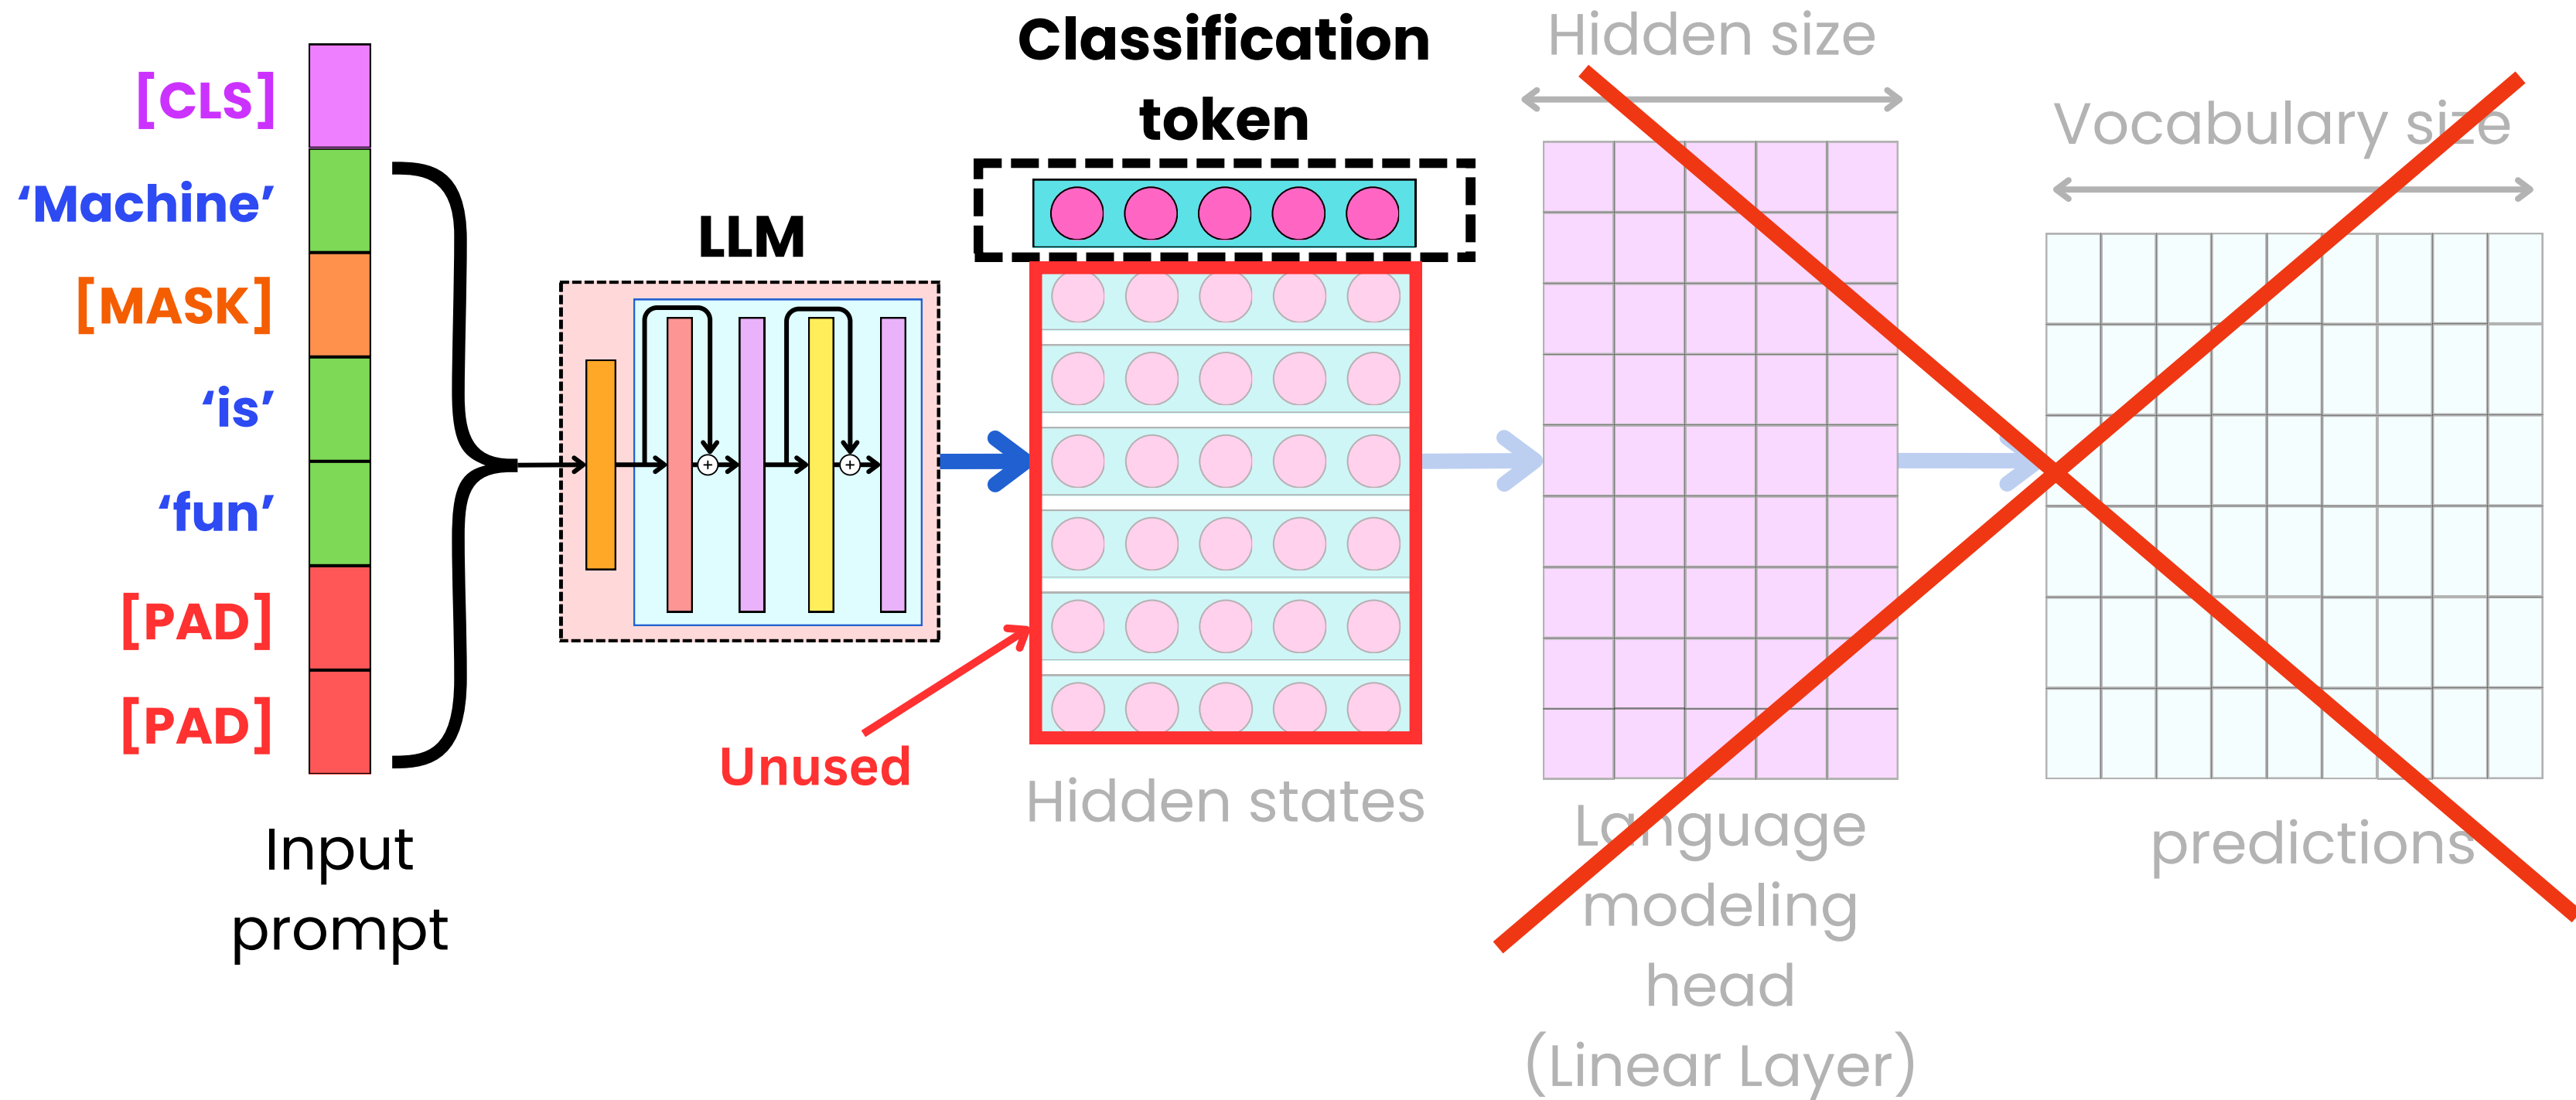


Masked Language Modeling









Masked Language Modeling




Sequence Prediction



Sequence Prediction




Sequence Prediction





Sequence Prediction


- Machine translation
- Text summarization
- Question-answering
- Multimodal modeling



Sequence-to-sequence




Sequence Prediction






Sequence Prediction






Text Classification






Text Classification



We start from a pre-trained LLM




Text Classification






Text Classification



We replace the language modeling head by a classification head




Text Classification



We only use ONE hidden state




Text Classification



**We can now compare the predictions
to the targets**




Token Classification






Token Classification






Token Classification






Token Classification



We replace the language modeling head by a classification head




Token Classification






Text Encoding






Text Encoding






Text Encoding






Text Encoding





Text Encoding

**“A DOG ON
THE GRASS”**



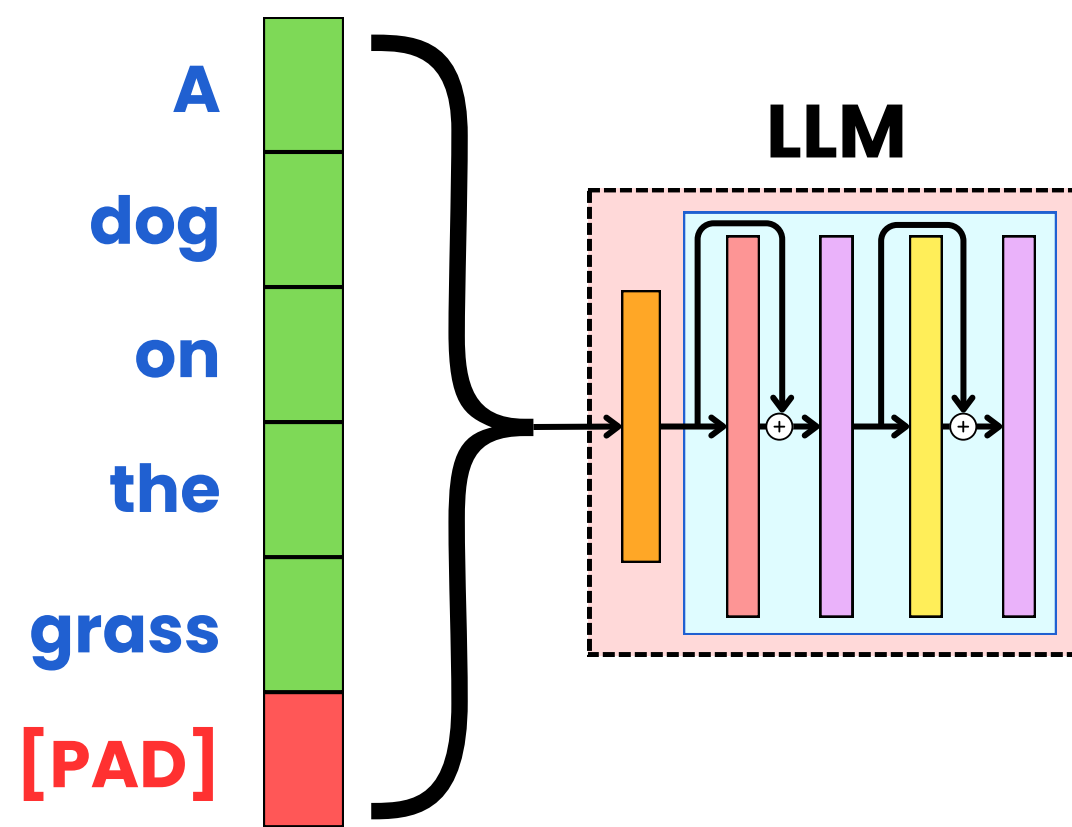

Vector database

**How do we
retrieve the
right images?**






Text Encoding

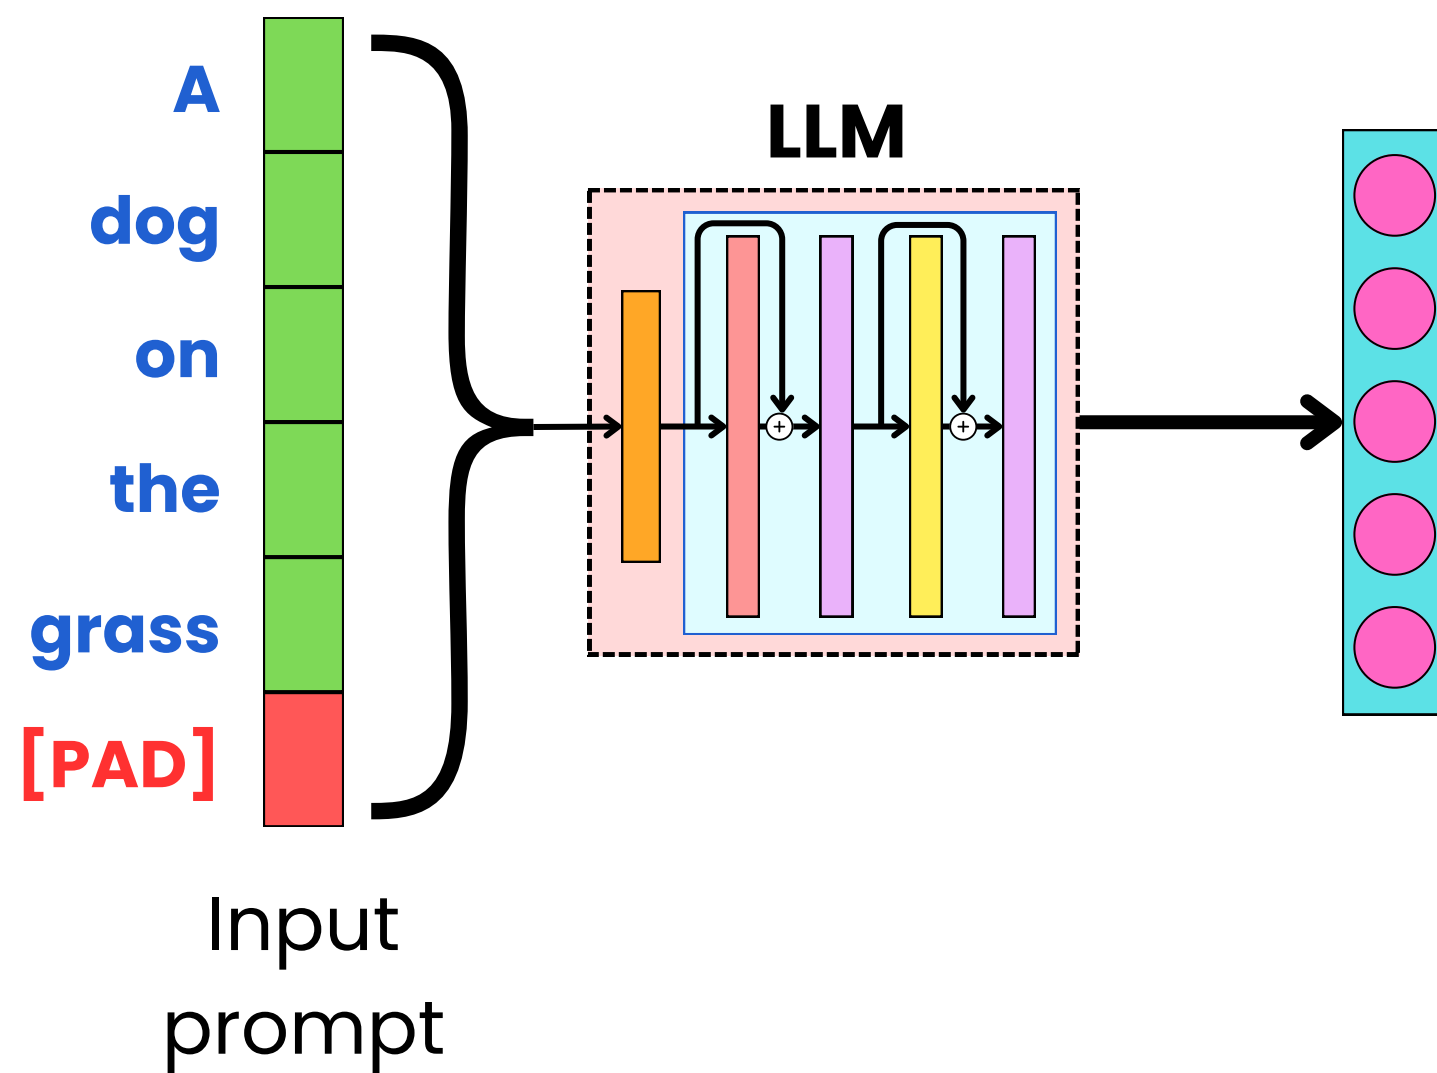



Input
prompt






Text Encoding






Text Encoding



**Train LLM to align
with Computer
Vision model**





Multimodal Fine-Tuning



Describe

the

image

[SEP]

A

dog

the

image

[SEP]

A

dog

[PAD]


Labels

Input data




Multimodal Fine-Tuning

Image encoder



**Vision-Language
connector**

Visual tokens




Text tokens

**Describe
the
image
[SEP]
A
dog**

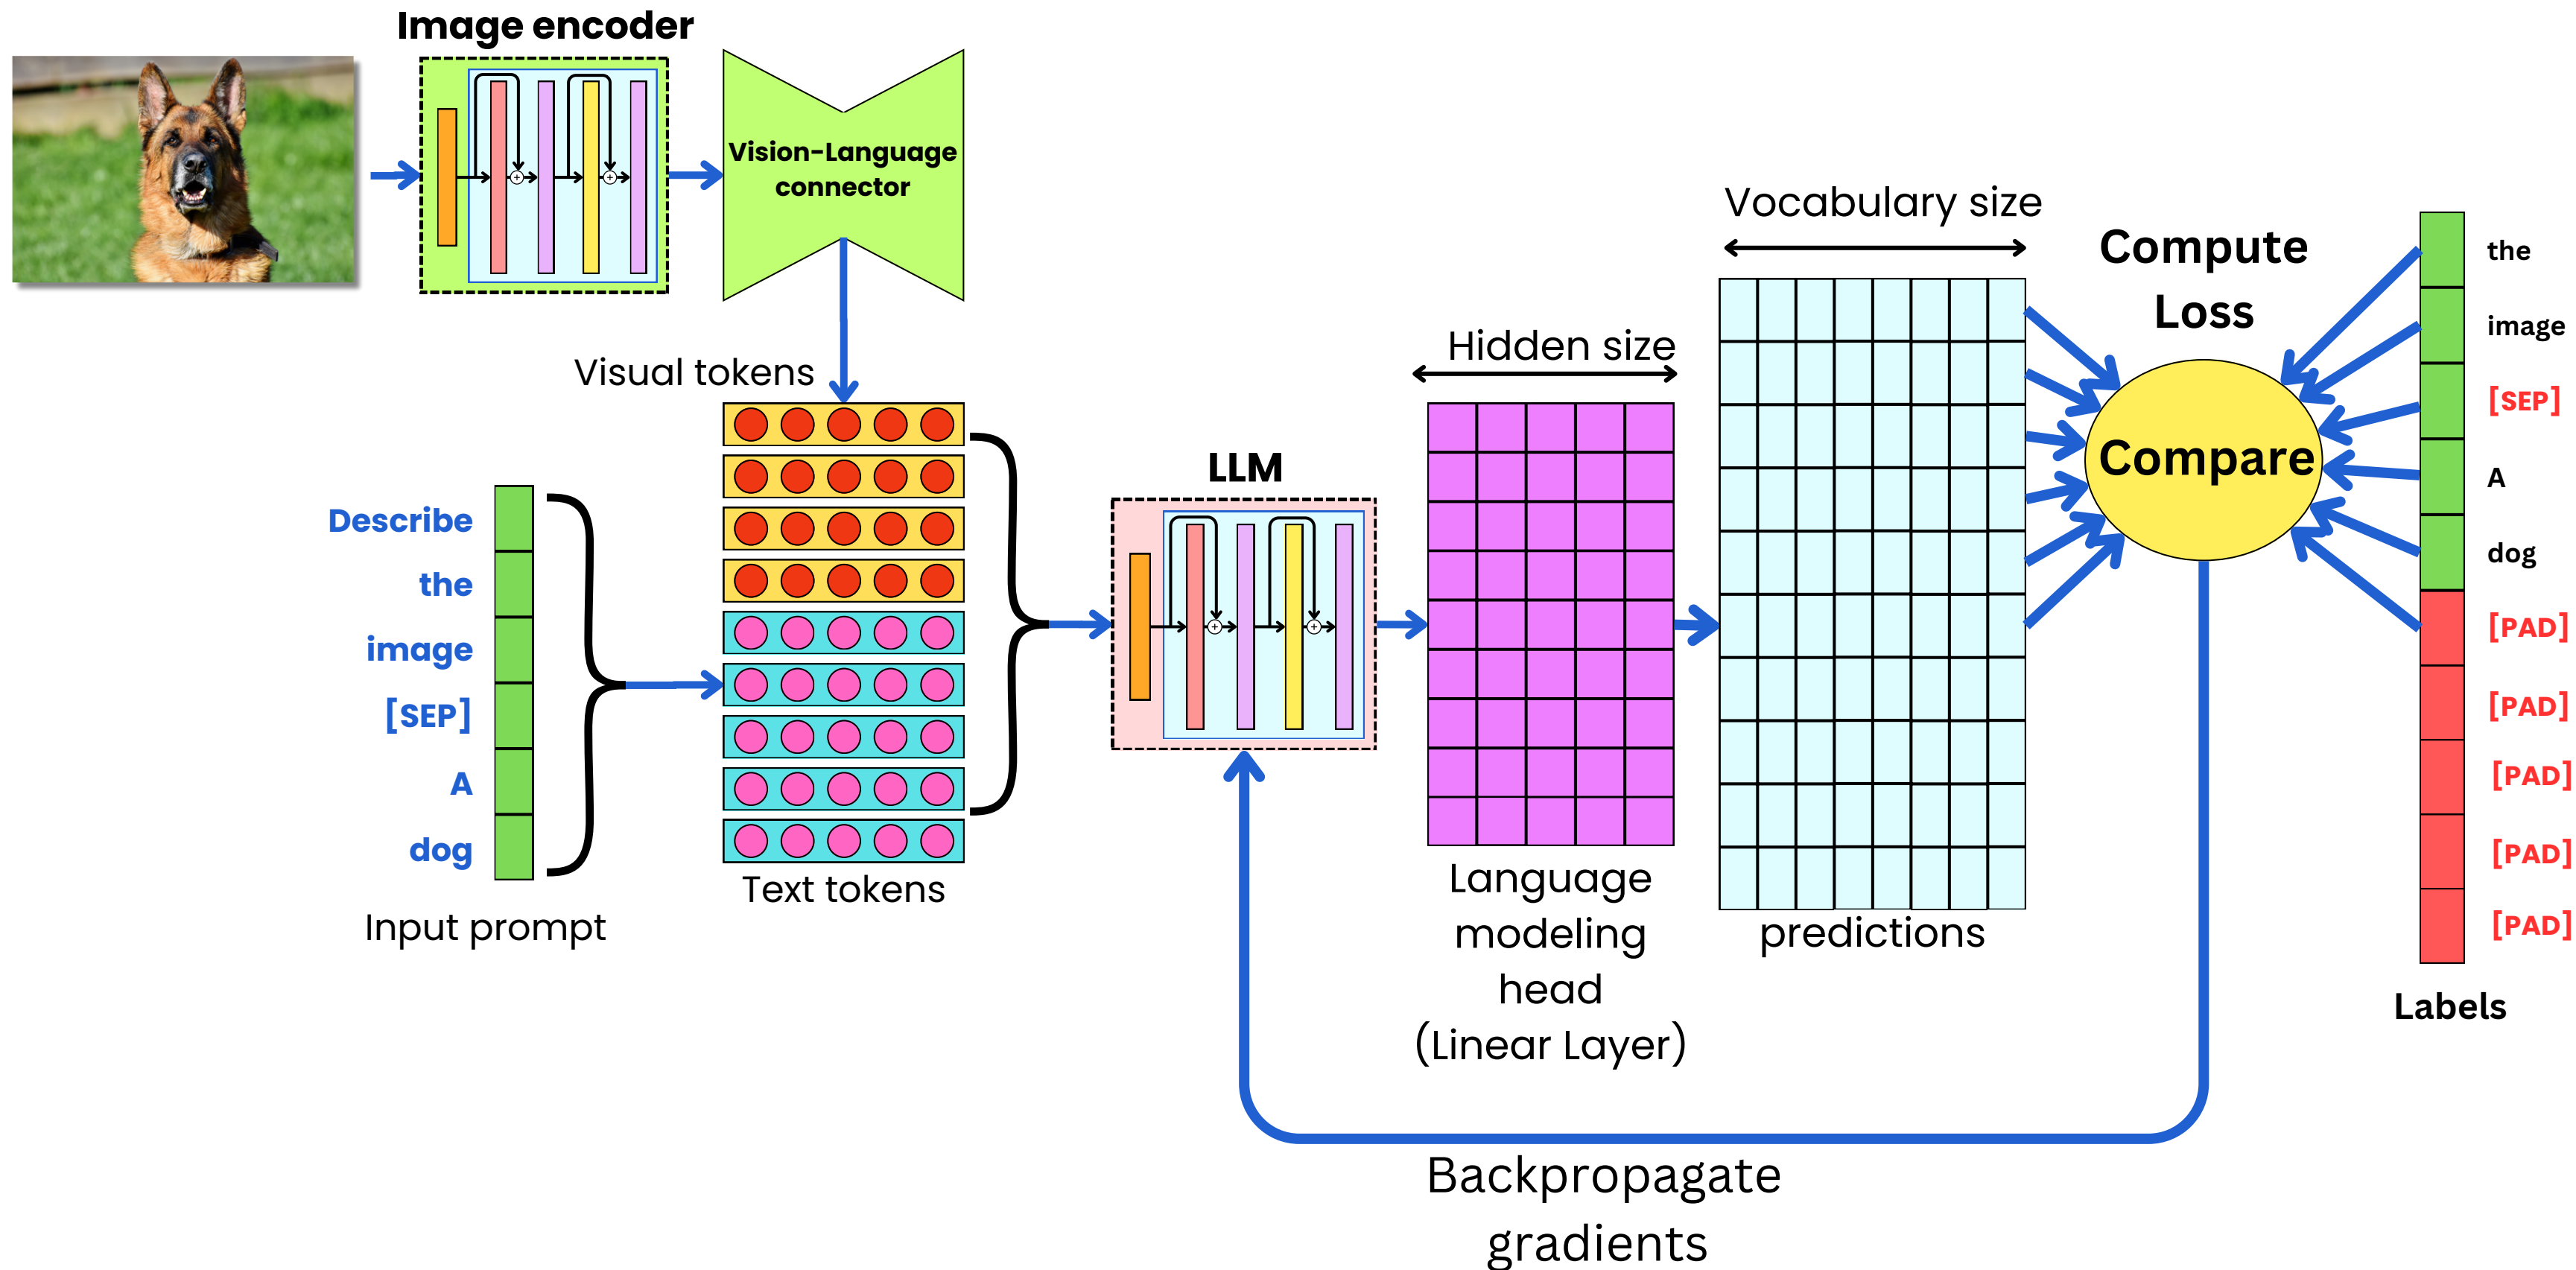
Input prompt

LLM



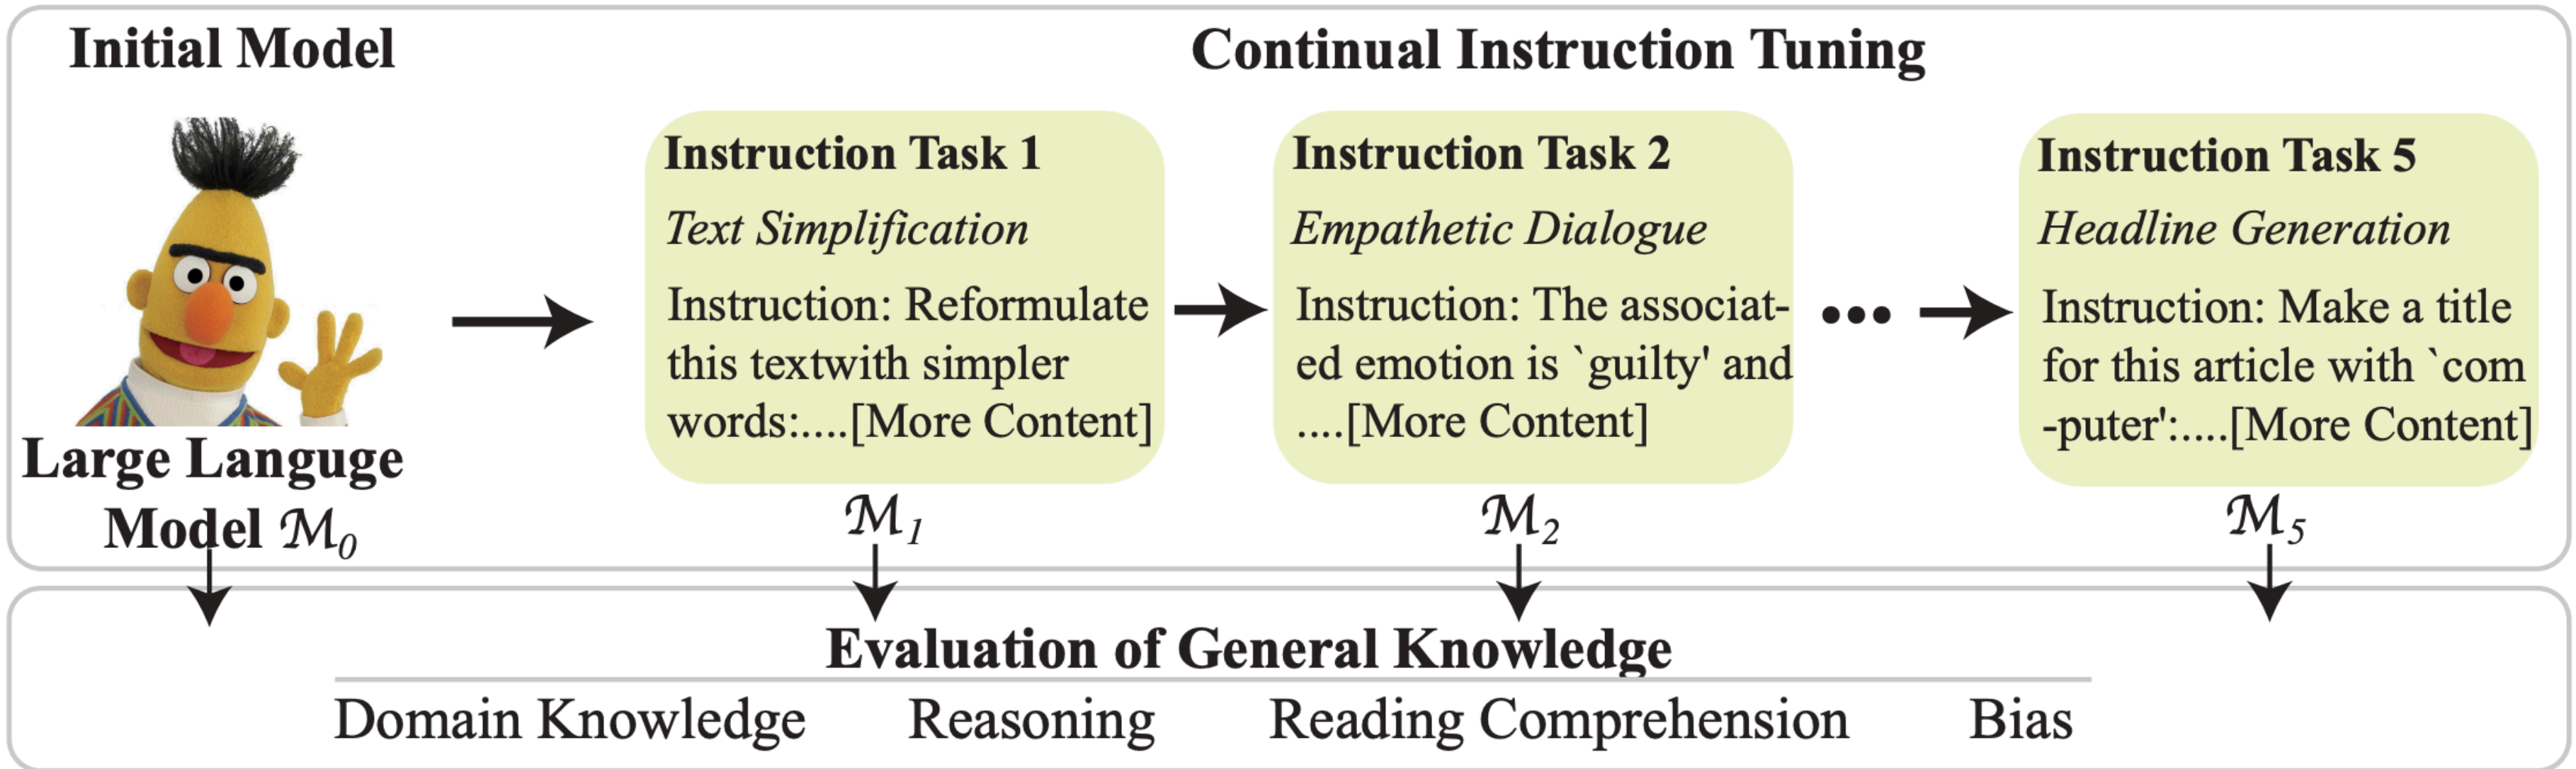


Multimodal Fine-Tuning






Catastrophic Forgetting



An Empirical Study of Catastrophic Forgetting in Large
Language Models During Continual Fine-tuning
<https://arxiv.org/pdf/2308.08747>



Catastrophic Forgetting



An Empirical Study of Catastrophic Forgetting in Large
Language Models During Continual Fine-tuning
<https://arxiv.org/pdf/2308.08747>



LoRA: Low-Rank Adaptation

**The Gradient
descent update**

$$\theta_t \leftarrow \theta_{t-1} - \alpha \nabla_{\theta} \mathcal{L} \big|_{\theta=\theta_{t-1}}$$

A trained model

$$\theta_T = \theta_0 - \alpha \nabla_{\theta} \mathcal{L} \big|_{\theta=\theta_0} - \alpha \nabla_{\theta} \mathcal{L} \big|_{\theta=\theta_1} \dots - \alpha \nabla_{\theta} \mathcal{L} \big|_{\theta=\theta_{T-1}}$$



LoRA: Low-Rank Adaptation

The Gradient
descent update

$$\theta_t \leftarrow \theta_{t-1} - \alpha \nabla_{\theta} \mathcal{L} |_{\theta=\theta_{t-1}}$$

A trained model

$$\theta_T = \theta_0 - \alpha \nabla_{\theta} \mathcal{L} |_{\theta=\theta_0} - \alpha \nabla_{\theta} \mathcal{L} |_{\theta=\theta_1} \dots - \alpha \nabla_{\theta} \mathcal{L} |_{\theta=\theta_{T-1}}$$

A fine-tuned
model

$$\theta_F = \theta_0 - \alpha \sum_{t=0}^{T-1} \nabla_{\theta} \mathcal{L} |_{\theta=\theta_t} - \alpha \sum_{t=T}^{T+F} \nabla_{\theta} \mathcal{L} |_{\theta=\theta_t}$$

The trained model

The fine-tuning data



LoRA: Low-Rank Adaptation

The Gradient
descent update

$$\theta_t \leftarrow \theta_{t-1} - \alpha \nabla_{\theta} \mathcal{L} |_{\theta=\theta_{t-1}}$$

A trained model

$$\theta_T = \theta_0 - \alpha \nabla_{\theta} \mathcal{L} |_{\theta=\theta_0} - \alpha \nabla_{\theta} \mathcal{L} |_{\theta=\theta_1} \dots - \alpha \nabla_{\theta} \mathcal{L} |_{\theta=\theta_{T-1}}$$

A fine-tuned
model

$$\theta_F = \theta_0 - \alpha \sum_{t=0}^{T-1} \nabla_{\theta} \mathcal{L} |_{\theta=\theta_t} - \alpha \sum_{t=T}^{T+F} \nabla_{\theta} \mathcal{L} |_{\theta=\theta_t}$$

The trained model

The fine-tuning data

$$\theta_F = \theta_T + \Delta W$$

The trained model

The fine-tuning data




LoRA: Low-Rank Adaptation



$$\theta_F = \boxed{\theta_T} + \boxed{\Delta W}$$

The trained model

The fine-tuning data



Augment layers with
LoRA Adaptors






LoRA: Low-Rank Adaptation

ΔW

same dimension as

θ_T




low-rank approximation




LoRA: Low-Rank Adaptation

$$\theta_F = \theta_T + \Delta W$$


The trained model The fine-tuning data



Augment layers with LoRA Adaptors



Append new matrix




Append new matrix




LoRA: Low-Rank Adaptation

$$\theta_F = \underbrace{\theta_T}_{\text{The trained model}} + \underbrace{\Delta W}_{\text{The fine-tuning data}} \simeq \theta_T + \underbrace{B^T A}_{\text{Low-Rank approximation}}$$



Augment layers with
LoRA Adaptors




Forward pass



LoRA: Low-Rank Adaptation

$$A_t \leftarrow A_{t-1} - \alpha \nabla_A \mathcal{L} \big|_{A=A_{t-1}}$$


$$B_t \leftarrow B_{t-1} - \alpha \nabla_B \mathcal{L} \big|_{B=B_{t-1}}$$



Backward pass



LoRA: Low-Rank Adaptation



Number of operations for Ax^T

$$1 \times R \times r = 10,000 \times 4 = 40,000$$

Number of operations for $B(Ax^T)$

$$1 \times r \times C = 4 \times 20,000 = 80,000$$

Total Number of operations

$$40,000 + 80,000 = 120,000$$

Number of elements

$$40,000 + 80,000 = 120,000$$

Number of operations

$$1 \times R \times C = 10,000 \times 20,000 = 200,000,000$$

Number of elements

$$1 \times R \times C = 10,000 \times 20,000 = 200,000,000$$




Precision

GPU

Model parameters (float16)



Precision



Adam

Momentum

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) \nabla_{\theta} \mathcal{L}$$

Variance


$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) (\nabla_{\theta} \mathcal{L})^2$$

Weights update

$$\theta_t = \theta_{t-1} - \frac{\mu}{\sqrt{\hat{v}_t + \epsilon}} \hat{m}_t$$



Precision



Adam

Momentum

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) \nabla_{\theta} \mathcal{L}$$

Variance

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) (\nabla_{\theta} \mathcal{L})^2$$

Weights update

$$\theta_t = \theta_{t-1} - \frac{\mu}{\sqrt{\hat{v}_t + \epsilon}} \hat{m}_t$$



Precision

GPU

Model parameters (float16)



Precision

Float 32

8 bits

23 bits



Sign ← Exponent × Mantissa == decimal →

$$(-1)^{sign} 2^{(exponent-127)} \times 1.mantissa$$

1 → 1

01010001 → 81

1 01100001000110100000110 → 1.37930369377136230469


$$(-1)^1 2^{(81-127)} \times 1.37930369377136230469$$

$$= -1.9601084e-14$$



Precision

Float 16



Sign ← Exponent → Mantissa == decimal →

$$(-1)^{sign} 2^{(exponent-15)} \times 1.mantissa$$

Brain Float 16



Sign ← Exponent → Mantissa →

$$(-1)^{sign} 2^{(exponent-127)} \times 1.mantissa$$

Float 8



Sign Exponent Mantissa

$$(-1)^{sign} 2^{(exponent-7)} \times 1.mantissa$$



Precision

Float 32 $\left[-3.4 \times 10^{38}, 3.4 \times 10^{38}\right]$

Float 16 $\left[-6.55 \times 10^4, 6.55 \times 10^4\right]$


BFloat 16 $\left[-3.39 \times 10^{38}, 3.39 \times 10^{38}\right]$

Float 8 $\left[-240, 240\right]$



Converting Data Type

Float 32



Float 16



Can float overflow



Converting Data Type

Float 32

1	0	1	0	1	0	0	0	1	0	1	1	0	0	0	0	1	0	0	0	1	1	0	1	0	0	0	0	0	1	1	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---




BFloat 16

1	0	1	0	1	0	0	0	1	0	1	1	0	0	0	0	1	0	0	0	1	1	0	1	0	0	0	0	0	0	1	1	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---


Does not overflow



Quantization



Converting to lower
precision data type




Information loss!




Quantization

Original data type




Rescale to unit range by using the data available






Quantization


Original data type



Rescale to unit range by using the data available




Rescale to the data type target range (e.g. Int8)






Quantization

Original data type




**Adaptive data type
conversion**


Rescale to unit range by using the data available



Rescale to the data type target range (e.g. Int8)




Convert to target data type (e.g. Int8)






Quantization to 4-bits NormalFloat




BFloat 16






Quantization to 4-bits NormalFloat

BFloat 16



Naive 4-bits quantization




**Many data points end up in
the same bucket**




Quantization to 4-bits NormalFloat

BFloat 16




Normal distributed 4-bits quantization




Less lossy




QLoRA









QLoRA






QLoRA








QLoRA







- 4-bits NormalFloat quantization
- Double quantization (for the quantization constants)
- Paged Optimizer